

e8uoزالak

December 23, 2024

LoanTap is an online platform committed to delivering customized loan products to millennials. They innovate in an otherwise dull loan segment, to deliver instant, flexible loans on consumer friendly terms to salaried professionals and businessmen.

The data science team at LoanTap is building an underwriting layer to determine the creditworthiness of MSMEs as well as individuals.

LoanTap deploys formal credit to salaried individuals and businesses 4 main financial instruments:
* Personal Loan * EMI Free Loan * Personal Overdraft * Advance Salary Loan

This case study will focus on the underwriting process behind Personal Loan only

Given a set of attributes for an Individual, determine if a credit line should be extended to them. If so, what should the repayment terms be in business recommendations?

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: !wget 1ZPYj7CZCfxntE8p2Lze_4Q04MyEOy6_d
```

Downloading...

From: https://drive.google.com/uc?id=1ZPYj7CZCfxntE8p2Lze_4Q04MyEOy6_d

To: /content/logistic_regression.csv

100% 100M/100M [00:01<00:00, 87.2MB/s]

```
[3]: df=pd.read_csv('logistic_regression.csv')
```

```
[ ]: df
```

```
[ ]:
      loan_amnt      term  int_rate  installment  grade  sub_grade  \
0      10000.0    36 months    11.44         329.48     B         B4
1       8000.0    36 months    11.99         265.68     B         B5
2      15600.0    36 months    10.49         506.97     B         B3
3       7200.0    36 months     6.49         220.65     A         A2
4      24375.0    60 months    17.27         609.33     C         C5
...      ...      ...      ...      ...      ...      ...
396025  10000.0    60 months    10.99         217.38     B         B4
396026  21000.0    36 months    12.29         700.42     C         C1
396027   5000.0    36 months     9.99         161.32     B         B1
396028  21000.0    60 months    15.31         503.02     C         C2
```

396029	2000.0	36 months	13.61	67.98	C	C2
--------	--------	-----------	-------	-------	---	----

	emp_title	emp_length	home_ownership	annual_inc	...	\
0	Marketing	10+ years	RENT	117000.0	...	
1	Credit analyst	4 years	MORTGAGE	65000.0	...	
2	Statistician	< 1 year	RENT	43057.0	...	
3	Client Advocate	6 years	RENT	54000.0	...	
4	Destiny Management Inc.	9 years	MORTGAGE	55000.0	...	
...	
396025	licensed bankere	2 years	RENT	40000.0	...	
396026	Agent	5 years	MORTGAGE	110000.0	...	
396027	City Carrier	10+ years	RENT	56500.0	...	
396028	Gracon Services, Inc	10+ years	MORTGAGE	64000.0	...	
396029	Internal Revenue Service	10+ years	RENT	42996.0	...	

	open_acc	pub_rec	revol_bal	revol_util	total_acc	initial_list_status	\
0	16.0	0.0	36369.0	41.8	25.0		w
1	17.0	0.0	20131.0	53.3	27.0		f
2	13.0	0.0	11987.0	92.2	26.0		f
3	6.0	0.0	5472.0	21.5	13.0		f
4	13.0	0.0	24584.0	69.8	43.0		f
...	
396025	6.0	0.0	1990.0	34.3	23.0		w
396026	6.0	0.0	43263.0	95.7	8.0		f
396027	15.0	0.0	32704.0	66.9	23.0		f
396028	9.0	0.0	15704.0	53.8	20.0		f
396029	3.0	0.0	4292.0	91.3	19.0		f

	application_type	mort_acc	pub_rec_bankruptcies	\
0	INDIVIDUAL	0.0	0.0	
1	INDIVIDUAL	3.0	0.0	
2	INDIVIDUAL	0.0	0.0	
3	INDIVIDUAL	0.0	0.0	
4	INDIVIDUAL	1.0	0.0	
...	
396025	INDIVIDUAL	0.0	0.0	
396026	INDIVIDUAL	1.0	0.0	
396027	INDIVIDUAL	0.0	0.0	
396028	INDIVIDUAL	5.0	0.0	
396029	INDIVIDUAL	NaN	0.0	

	address
0	0174 Michelle Gateway\r\nMendozaberg, OK 22690
1	1076 Carney Fort Apt. 347\r\nLoganmouth, SD 05113
2	87025 Mark Dale Apt. 269\r\nNew Sabrina, WV 05113
3	823 Reid Ford\r\nDelacruzside, MA 00813
4	679 Luna Roads\r\nGreggshire, VA 11650

```

...
396025 12951 Williams Crossing\r\nJohnnyville, DC 30723
396026 0114 Fowler Field Suite 028\r\nRachelborough, ...
396027 953 Matthew Points Suite 414\r\nReedfort, NY 7...
396028 7843 Blake Freeway Apt. 229\r\nNew Michael, FL...
396029 787 Michelle Causeway\r\nBriannaton, AR 48052

```

[396030 rows x 27 columns]

Data Description

1. loan_amnt : The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.
2. term : The number of payments on the loan. Values are in months and can be either 36 or 60.
3. int_rate : Interest Rate on the loan
4. installment : The monthly payment owed by the borrower if the loan originates.
5. grade : LoanTap assigned loan grade
6. sub_grade : LoanTap assigned loan subgrade
7. emp_title : The job title supplied by the Borrower when applying for the loan.
8. emp_length : Employment length in years. Possible values are between 0 and 10 where 0 means less than one year and 10 means ten or more years.
9. home_ownership : The home ownership status provided by the borrower during registration or obtained from the credit report.
10. annual_inc : The self-reported annual income provided by the borrower during registration.
11. verification_status : Indicates if income was verified by LoanTap, not verified, or if the income source was verified
12. issue_d : The month which the loan was funded
13. loan_status : Current status of the loan - Target Variable
14. purpose : A category provided by the borrower for the loan request.
15. title : The loan title provided by the borrower
16. dti : A ratio calculated using the borrower's total monthly debt payments on the total debt obligations, excluding mortgage and the requested LoanTap loan, divided by the borrower's self-reported monthly income.
17. earliest_cr_line : The month the borrower's earliest reported credit line was opened
18. open_acc : The number of open credit lines in the borrower's credit file.
19. pub_rec : Number of derogatory public records
20. revol_bal : Total credit revolving balance
21. revol_util : Revolving line utilization rate, or the amount of credit the borrower is using relative to all available revolving credit.
22. total_acc : The total number of credit lines currently in the borrower's credit file
23. initial_list_status : The initial listing status of the loan. Possible values are – W, F
24. application_type : Indicates whether the loan is an individual application or a joint application with two co-borrowers
25. mort_acc : Number of mortgage accounts.
26. pub_rec_bankruptcies : Number of public record bankruptcies
27. Address: Address of the individual

###1.Define Problem Statement and perform Exploratory Data Analysis

1.Definition of problem

The motive of this study is that from given set of attributes for an Individual, determine if a credit line should be extended to them. The attributes include loan statistics like loan amount, rate of interest ect., and applicant statistics like annual income, home ownership etc.

2.Observations on shape of data, data types of all the attributes, conversion of categorical attributes to 'category' (If required), missing value detection, statistical summary

```
[ ]: df.shape
```

```
[ ]: (396030, 27)
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             396030 non-null  float64
1   term                  396030 non-null  object
2   int_rate              396030 non-null  float64
3   installment           396030 non-null  float64
4   grade                 396030 non-null  object
5   sub_grade             396030 non-null  object
6   emp_title             373103 non-null  object
7   emp_length            377729 non-null  object
8   home_ownership        396030 non-null  object
9   annual_inc            396030 non-null  float64
10  verification_status   396030 non-null  object
11  issue_d               396030 non-null  object
12  loan_status           396030 non-null  object
13  purpose               396030 non-null  object
14  title                 394274 non-null  object
15  dti                   396030 non-null  float64
16  earliest_cr_line      396030 non-null  object
17  open_acc              396030 non-null  float64
18  pub_rec               396030 non-null  float64
19  revol_bal             396030 non-null  float64
20  revol_util            395754 non-null  float64
21  total_acc             396030 non-null  float64
22  initial_list_status   396030 non-null  object
23  application_type      396030 non-null  object
24  mort_acc              358235 non-null  float64
25  pub_rec_bankruptcies  395495 non-null  float64
```

```

26 address 396030 non-null object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB

```

```

[ ]: #missing value detection
df.isna().sum()

```

```

[ ]: loan_amnt      0
term              0
int_rate          0
installment       0
grade             0
sub_grade         0
emp_title        22927
emp_length       18301
home_ownership    0
annual_inc        0
verification_status 0
issue_d           0
loan_status       0
purpose           0
title            1756
dti               0
earliest_cr_line  0
open_acc          0
pub_rec           0
revol_bal         0
revol_util        276
total_acc         0
initial_list_status 0
application_type  0
mort_acc          37795
pub_rec_bankruptcies 535
address           0
dtype: int64

```

```

[ ]: #statistical summary of the dataset
df.describe()

```

```

[ ]:
count    loan_amnt    int_rate    installment    annual_inc  \
count  396030.000000  396030.000000  396030.000000  3.960300e+05
mean    14113.888089    13.639400    431.849698  7.420318e+04
std      8357.441341     4.472157    250.727790  6.163762e+04
min       500.000000     5.320000     16.080000  0.000000e+00
25%      8000.000000    10.490000    250.330000  4.500000e+04
50%     12000.000000    13.330000    375.430000  6.400000e+04
75%     20000.000000    16.490000    567.300000  9.000000e+04

```

max	40000.000000	30.990000	1533.810000	8.706582e+06
-----	--------------	-----------	-------------	--------------

	dti	open_acc	pub_rec	revol_bal \
count	396030.000000	396030.000000	396030.000000	3.960300e+05
mean	17.379514	11.311153	0.178191	1.584454e+04
std	18.019092	5.137649	0.530671	2.059184e+04
min	0.000000	0.000000	0.000000	0.000000e+00
25%	11.280000	8.000000	0.000000	6.025000e+03
50%	16.910000	10.000000	0.000000	1.118100e+04
75%	22.980000	14.000000	0.000000	1.962000e+04
max	9999.000000	90.000000	86.000000	1.743266e+06

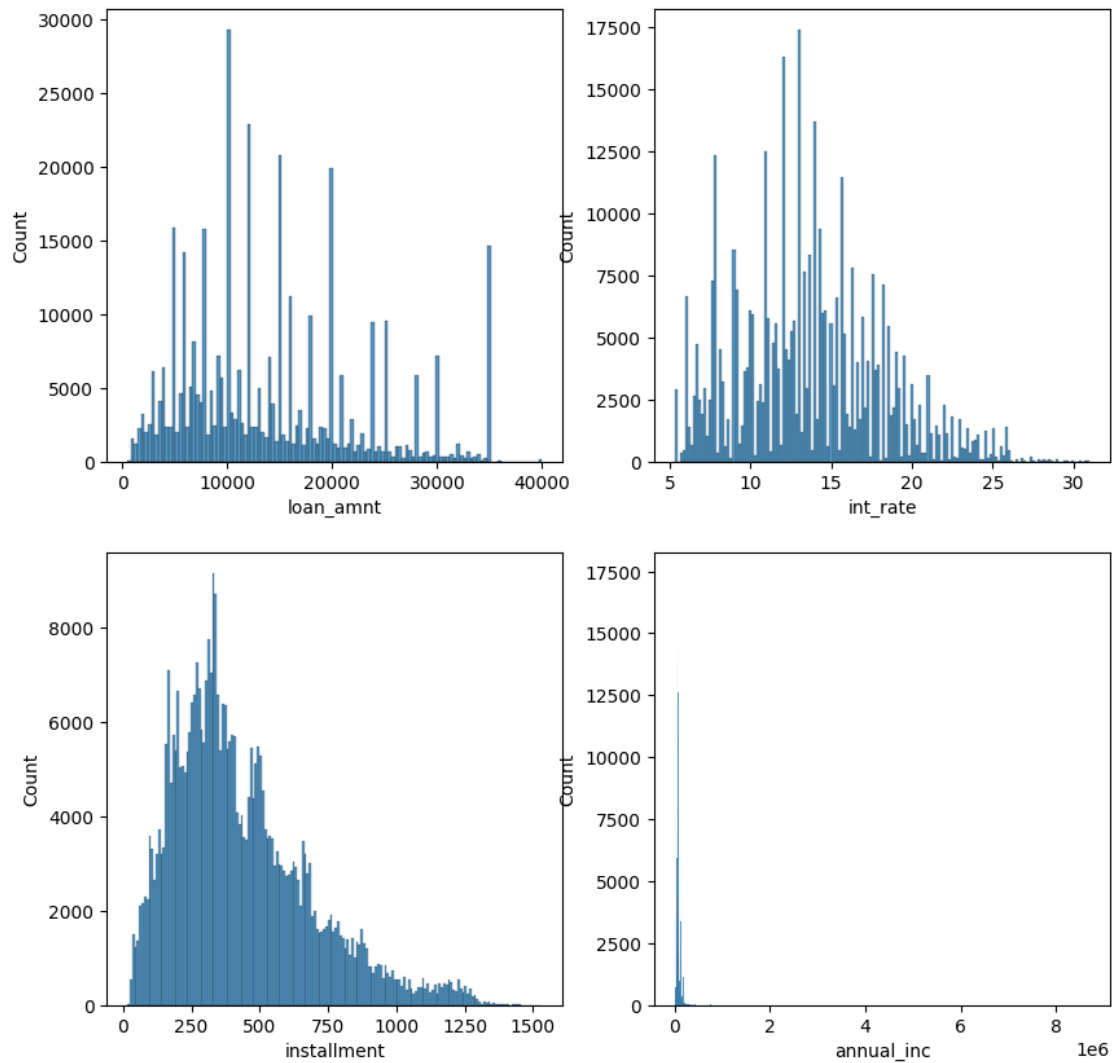
	revol_util	total_acc	mort_acc	pub_rec_bankruptcies
count	395754.000000	396030.000000	358235.000000	395495.000000
mean	53.791749	25.414744	1.813991	0.121648
std	24.452193	11.886991	2.147930	0.356174
min	0.000000	2.000000	0.000000	0.000000
25%	35.800000	17.000000	0.000000	0.000000
50%	54.800000	24.000000	1.000000	0.000000
75%	72.900000	32.000000	3.000000	0.000000
max	892.300000	151.000000	34.000000	8.000000

3.Univariate Analysis (distribution plots of all the continuous variable(s) barplots/countplots of all the categorical variables)

```
[ ]: import seaborn as sns
import matplotlib.pyplot as plt
```


```
[ ]: plt.figure(figsize=(10,10))
plt.subplot(2,2,1)
sns.histplot(df['loan_amnt'])
plt.subplot(2,2,2)
sns.histplot(df['int_rate'])
plt.subplot(2,2,3)
sns.histplot(df['installment'])
plt.subplot(2,2,4)
sns.histplot(df['annual_inc'])
```

```
[ ]: <Axes: xlabel='annual_inc', ylabel='Count'>
```



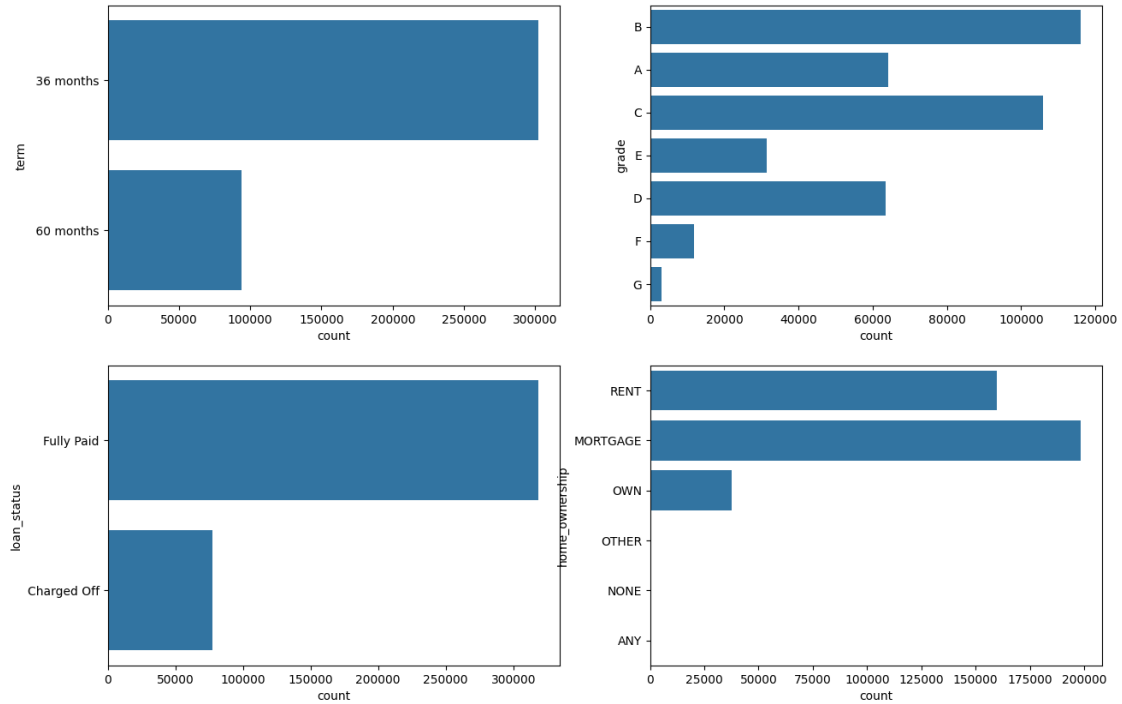
```
[ ]: sns.boxplot(df['annual_inc'])
```

```
[ ]: <Axes: ylabel='annual_inc'>
```



```
[ ]: #count plots for all categorical variables
plt.figure(figsize=(15,10))
plt.subplot(2,2,1)
sns.countplot(df['term'])
plt.subplot(2,2,2)
sns.countplot(df['grade'])
plt.subplot(2,2,3)
sns.countplot(df['loan_status'])
plt.subplot(2,2,4)
sns.countplot(df['home_ownership'])
```

```
[ ]: <Axes: xlabel='count', ylabel='home_ownership'>
```

4. Bivariate Analysis (Relationships between important variable)

```
[ ]: #correlation and heatmap of numerical fields
df.corr(numeric_only=True)
```

```
[ ]:
      loan_amnt  int_rate  installment  annual_inc  dti \
loan_amnt      1.000000  0.168921    0.953929    0.336887  0.016636
int_rate       0.168921  1.000000    0.162758   -0.056771  0.079038
installment    0.953929  0.162758    1.000000    0.330381  0.015786
annual_inc     0.336887 -0.056771    0.330381    1.000000 -0.081685
dti            0.016636  0.079038    0.015786   -0.081685  1.000000
open_acc       0.198556  0.011649    0.188973    0.136150  0.136181
pub_rec       -0.077779  0.060986   -0.067892   -0.013720 -0.017639
revol_bal      0.328320 -0.011280    0.316455    0.299773  0.063571
revol_util     0.099911  0.293659    0.123915    0.027871  0.088375
total_acc      0.223886 -0.036404    0.202430    0.193023  0.102128
mort_acc       0.222315 -0.082583    0.193694    0.236320 -0.025439
pub_rec_bankruptcies -0.106539  0.057450   -0.098628   -0.050162 -0.014558


      open_acc  pub_rec  revol_bal  revol_util  total_acc \
loan_amnt     0.198556 -0.077779    0.328320    0.099911    0.223886
int_rate      0.011649  0.060986   -0.011280    0.293659   -0.036404
installment   0.188973 -0.067892    0.316455    0.123915    0.202430
annual_inc    0.136150 -0.013720    0.299773    0.027871    0.193023
```

dti	0.136181	-0.017639	0.063571	0.088375	0.102128
open_acc	1.000000	-0.018392	0.221192	-0.131420	0.680728
pub_rec	-0.018392	1.000000	-0.101664	-0.075910	0.019723
revol_bal	0.221192	-0.101664	1.000000	0.226346	0.191616
revol_util	-0.131420	-0.075910	0.226346	1.000000	-0.104273
total_acc	0.680728	0.019723	0.191616	-0.104273	1.000000
mort_acc	0.109205	0.011552	0.194925	0.007514	0.381072
pub_rec_bankruptcies	-0.027732	0.699408	-0.124532	-0.086751	0.042035

	mort_acc	pub_rec_bankruptcies
loan_amnt	0.222315	-0.106539
int_rate	-0.082583	0.057450
installment	0.193694	-0.098628
annual_inc	0.236320	-0.050162
dti	-0.025439	-0.014558
open_acc	0.109205	-0.027732
pub_rec	0.011552	0.699408
revol_bal	0.194925	-0.124532
revol_util	0.007514	-0.086751
total_acc	0.381072	0.042035
mort_acc	1.000000	0.027239
pub_rec_bankruptcies	0.027239	1.000000

```
[ ]: plt.figure(figsize=(12,10))
sns.heatmap(df.corr(numeric_only=True),annot=True)
```

```
[ ]: <Axes: >
```



5. Illustrate the insights based on EDA

1. The dataset contains 396 thousand entries and 27 columns
2. The columns are 12 float fields and 15 object type fields.
3. 'Annual income' have low range and high number of outliers.
4. The target variable loan status is unbalanced.
5. Majority of the loans are fully paid, 20% of the total loans being charged off.
6. Majority of loans are tenured for 36 months. Remaining are tenured for 60 months.
7. Greatest number of loans are graded B and C. Followed by A and D.
8. Interest rate and loan amount have strong positive correlation.
9. Other fields with strong positive correlation are pub_rec and pub_rec_bankruptcies, and total_acc and open_acc.

###2. Data Preprocessing

1. Duplicate value check

```
[ ]: #duplicate check  
df.duplicated().sum()
```

```
[ ]: 0
```

2.Missing value treatment

```
[ ]: df.isna().sum()
```

```
[ ]: loan_amnt          0  
term                  0  
int_rate              0  
installment           0  
grade                 0  
sub_grade             0  
emp_title             22927  
emp_length            18301  
home_ownership         0  
annual_inc            0  
verification_status    0  
issue_d               0  
loan_status            0  
purpose                0  
title                 1756  
dti                    0  
earliest_cr_line       0  
open_acc              0  
pub_rec               0  
revol_bal              0  
revol_util            276  
total_acc             0  
initial_list_status     0  
application_type        0  
mort_acc              37795  
pub_rec_bankruptcies    535  
address                0  
dtype: int64
```

```
[ ]: #removing missing entries in relevant columns  
df.dropna(subset=['emp_length'],inplace=True)
```

```
[ ]: df.isna().sum()
```

```
[ ]: loan_amnt          0  
term                  0  
int_rate              0  
installment           0
```

grade	0
sub_grade	0
emp_title	4804
emp_length	0
home_ownership	0
annual_inc	0
verification_status	0
issue_d	0
loan_status	0
purpose	0
title	1542
dti	0
earliest_cr_line	0
open_acc	0
pub_rec	0
revol_bal	0
revol_util	265
total_acc	0
initial_list_status	0
application_type	0
mort_acc	36739
pub_rec_bankruptcies	535
address	0
dtype:	int64

```
[ ]: df.dropna(subset=['mort_acc', 'pub_rec_bankruptcies', 'revol_util'], inplace=True)
```

```
[ ]: df.isna().sum()
```

loan_amnt	0
term	0
int_rate	0
installment	0
grade	0
sub_grade	0
emp_title	3376
emp_length	0
home_ownership	0
annual_inc	0
verification_status	0
issue_d	0
loan_status	0
purpose	0
title	1532
dti	0
earliest_cr_line	0
open_acc	0

```
pub_rec          0
revol_bal        0
revol_util       0
total_acc        0
initial_list_status 0
application_type 0
mort_acc         0
pub_rec_bankruptcies 0
address          0
dtype: int64
```

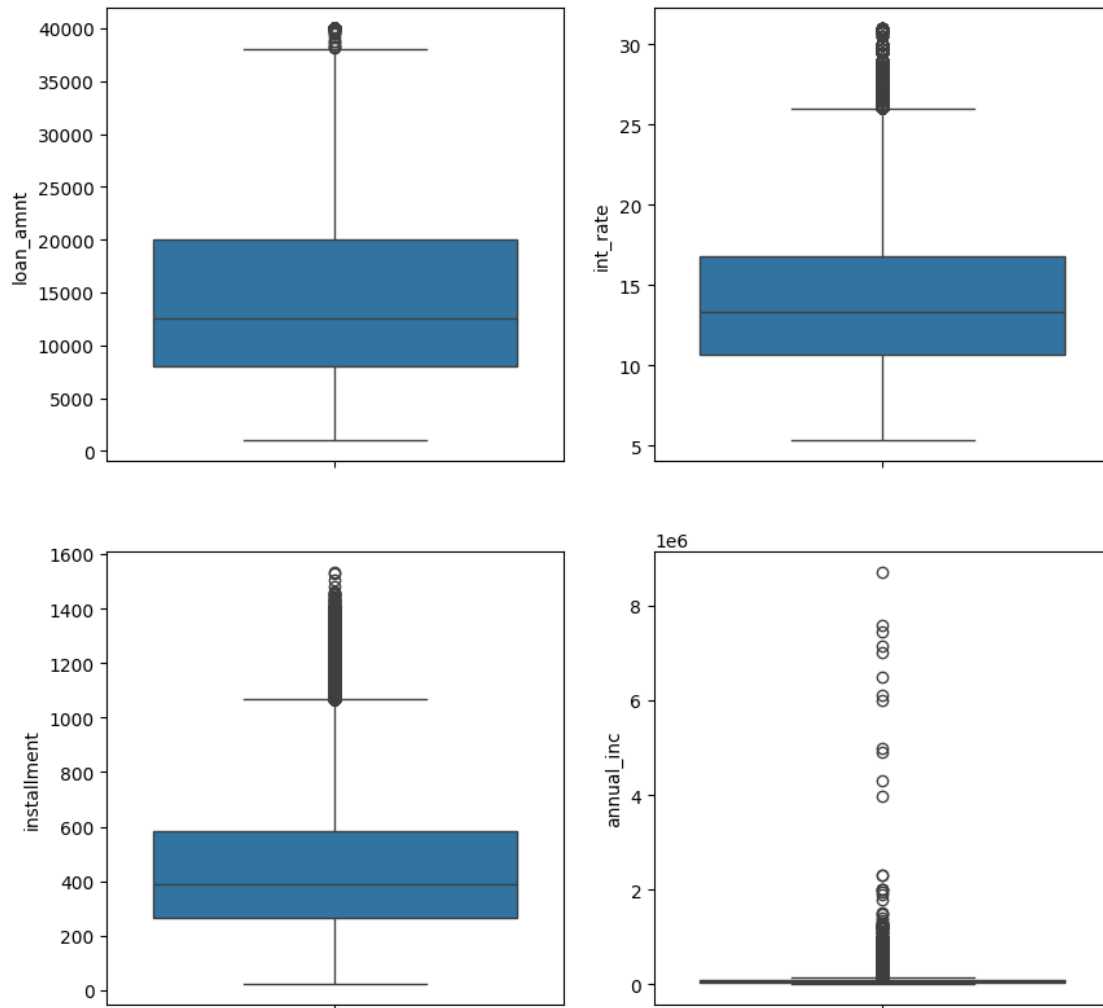
```
[ ]: df.shape
```

```
[ ]: (340775, 27)
```

3.Outlier treatment

```
[ ]: #boxplots of all continuous variables
plt.figure(figsize=(10,10))
plt.subplot(2,2,1)
sns.boxplot(df['loan_amnt'])
plt.subplot(2,2,2)
sns.boxplot(df['int_rate'])
plt.subplot(2,2,3)
sns.boxplot(df['installment'])
plt.subplot(2,2,4)
sns.boxplot(df['annual_inc'])
```

```
[ ]: <Axes: ylabel='annual_inc'>
```



4.Feature engineering

```
[ ]: #dropping irrelevant data for regression
df.
    ↳drop(columns=['emp_title','title','issue_d','earliest_cr_line','address'],inplace=True)

[ ]: df.drop(columns=['purpose'],inplace=True)

[ ]: #Creation of Flags- If value greater than 1.0 then 1 else 0. This can be done_
    ↳on: Pub_rec, Mort_acc and Pub_rec_bankruptcies
def flag(x):
    if x>1.0:
        return 1
    else:
        return 0
df['pub_rec_flag']=df['pub_rec'].apply(flag)
```

```
df['mort_acc_flag']=df['mort_acc'].apply(flag)
df['pub_rec_bankruptcies_flag']=df['pub_rec_bankruptcies'].apply(flag)
```

```
[ ]: #dropping original columns
df.drop(columns=['pub_rec','mort_acc','pub_rec_bankruptcies'],inplace=True)
```

```
[ ]: #converting term column into numerical
def term_num(x):
    if x==' 36 months':
        return 36
    else:
        return 60
df['term']=df['term'].apply(term_num)
```

```
[ ]: #converting emp length into numerical
def emp_length_num(x):
    if x=='< 1 year':
        return 0
    elif x=='1 year':
        return 1
    elif x=='2 years':
        return 2
    elif x=='3 years':
        return 3
    elif x=='4 years':
        return 4
    elif x=='5 years':
        return 5
    elif x=='6 years':
        return 6
    elif x=='7 years':
        return 7
    elif x=='8 years':
        return 8
    elif x=='9 years':
        return 9
    elif x=='10+ years':
        return 10
    else:
        return 0
df['emp_length']=df['emp_length'].apply(emp_length_num)
```

5.Data preparation for modeling

```
[ ]: #label for grade, subgrade, emp length, home ownership,application type,□
    ↪initial listing status etc
from sklearn.preprocessing import LabelEncoder
```



```

column = [
    'grade', 'sub_grade', 'home_ownership', 'application_type', 'initial_list_status', 'verification_status'
]
label_encoder = LabelEncoder()
for col in column:
    df[col] = label_encoder.fit_transform(df[col])

```

```
[ ]: df
```

```
[ ]:
      loan_amnt  term  int_rate  installment  grade  sub_grade  emp_length \
0      10000.0   36    11.44      329.48      1      8          10
1       8000.0   36    11.99      265.68      1      9           4
2      15600.0   36    10.49      506.97      1      7           0
3       7200.0   36     6.49      220.65      0      1           6
4      24375.0   60    17.27      609.33      2     14           9
...
396024      6000.0   36    13.11      202.49      1      8           5
396025     10000.0   60    10.99      217.38      1      8           2
396026     21000.0   36    12.29      700.42      2     10           5
396027      5000.0   36     9.99      161.32      1      5          10
396028     21000.0   60    15.31      503.02      2     11          10

```

```

      home_ownership  annual_inc  verification_status  ...  dti  open_acc \
0                5    117000.0                0  ...  26.24    16.0
1                1     65000.0                0  ...  22.05    17.0
2                5    43057.0                1  ...  12.79    13.0
3                5    54000.0                0  ...   2.60     6.0
4                1    55000.0                2  ...  33.95    13.0
...
396024            ...            ...            ...  ...  ...
396024                5     64000.0                0  ...  10.81     7.0
396025                5     40000.0                1  ...  15.63     6.0
396026                1    110000.0                1  ...  21.45     6.0
396027                5     56500.0                2  ...  17.56    15.0
396028                1     64000.0                2  ...  15.88     9.0

```

```

      revol_bal  revol_util  total_acc  initial_list_status \
0      36369.0      41.8      25.0                1
1      20131.0      53.3      27.0                0
2      11987.0      92.2      26.0                0
3       5472.0      21.5      13.0                0
4      24584.0      69.8      43.0                0
...
396024      11456.0      97.1       9.0                1
396025       1990.0      34.3      23.0                1
396026      43263.0      95.7       8.0                0
396027      32704.0      66.9      23.0                0
396028      15704.0      53.8      20.0                0

```

	application_type	pub_rec_flag	mort_acc_flag	\
0	1	0	0	
1	1	0	1	
2	1	0	0	
3	1	0	0	
4	1	0	0	
...	
396024	1	0	0	
396025	1	0	0	
396026	1	0	0	
396027	1	0	0	
396028	1	0	1	

	pub_rec_bankruptcies_flag
0	0
1	0
2	0
3	0
4	0
...	...
396024	0
396025	0
396026	0
396027	0
396028	0

[340775 rows x 21 columns]

```
[ ]: X=df.drop(columns=['loan_status'])
      y=df['loan_status']
```

```
[ ]: #coding y values into classes
      def class_code(x):
          if x=='Fully Paid':
              return 1
          else:
              return 0
      y=y.apply(class_code)
```

```
[ ]: #standard scaling
      from sklearn.preprocessing import StandardScaler
      scaler=StandardScaler()
      X = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
```

```
[ ]: #train-test split
      from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
↳2,random_state=42)
```

```
[ ]: X_train
```

```
[ ]:      loan_amnt      term  int_rate  installment      grade  sub_grade  \
179376    0.114022 -0.562923  0.113186     0.341142  0.113533   0.271920
337782    2.435052  1.776441  1.206435     1.848621  1.620561   1.488808
121817   -1.230985 -0.562923 -0.246056    -1.208891  0.113533  -0.032302
1642     -0.302573 -0.562923 -0.312582    -0.178983  0.113533  -0.184413
84796    -1.314304 -0.562923 -1.370349    -1.334970 -1.393495  -1.401302
...
119879   -0.296622 -0.562923 -0.623038    -0.204222 -0.639981  -0.488636
259178    1.697084 -0.562923 -0.033171     2.113617  0.113533  -0.184413
131932   -0.183546  1.776441 -1.086504    -0.700375 -1.393495  -1.097080
146867   -0.540628 -0.562923 -0.401283    -0.451269 -0.639981  -0.640747
121958   -0.969126 -0.562923 -0.179529    -0.912792 -0.639981  -0.336524

      emp_length  home_ownership  annual_inc  verification_status      dti  \
179376   -0.283824      0.586501    5.216758          -0.023409 -1.301416
337782    1.089033     -0.972102    1.352376          -0.023409 -0.283543
121817   -1.107539     -0.972102   -0.451002          -1.259991  1.895686
1642     1.089033     -0.972102    0.225265          -0.023409  0.601884
84796    -0.558396     -0.972102   -0.257783          -1.259991 -0.438063
...
119879    0.265319     -0.972102   -0.579815          -1.259991  1.115726
259178   -0.832967     1.106035    0.515093          -0.023409  0.518492
131932    1.089033     -0.972102   -0.483205           1.213173 -1.980815
146867    1.089033      0.586501   -0.096767          -1.259991 -0.365709
121958   -0.558396     -0.972102   -0.338291          -0.023409  1.737486

      open_acc  revol_bal  revol_util  total_acc  initial_list_status  \
179376  -1.078456   0.962220   -2.267513   -0.658432          -0.888529
337782  -0.113062   0.801693    0.740257    0.518216           1.125456
121817   1.431569  -0.506071   -1.132820    0.097985          -0.888529
1642     0.273096   0.037514    1.428582   -0.154154           1.125456
84796    0.466175  -0.046958    0.010216    1.526772          -0.888529
...
119879  -0.306141   0.133007    0.923810   -0.070108           1.125456
259178   0.466175   0.661457    1.553732    0.097985           1.125456
131932  -0.499220  -0.716590   -1.992183   -0.322247           1.125456
146867   0.659253  -0.382342   -0.169166   -0.154154           1.125456
121958   1.431569   0.146392   -0.123277   -0.406293          -0.888529

      application_type  pub_rec_flag  mort_acc_flag  \
179376          -0.007806        -0.1484      1.130197
337782          -0.007806        -0.1484      1.130197
```

121817	-0.007806	-0.1484	1.130197
1642	-0.007806	-0.1484	-0.884802
84796	-0.007806	-0.1484	1.130197
...
119879	-0.007806	-0.1484	1.130197
259178	-0.007806	-0.1484	-0.884802
131932	-0.007806	-0.1484	1.130197
146867	-0.007806	-0.1484	-0.884802
121958	-0.007806	-0.1484	-0.884802

	pub_rec_bankruptcies_flag
179376	-0.079943
337782	-0.079943
121817	-0.079943
1642	-0.079943
84796	-0.079943
...	...
119879	-0.079943
259178	-0.079943
131932	-0.079943
146867	-0.079943
121958	-0.079943

[272620 rows x 20 columns]

###3.Model building

1.Build the Logistic Regression model and comment on the model statistics

```
[ ]: #importing Logistic Regression
from sklearn.linear_model import LogisticRegression
logreg=LogisticRegression()
```

```
[ ]: #fitting the model
logreg.fit(X_train,y_train)
```

```
[ ]: LogisticRegression()
```

```
[ ]: #model accuracy
print('Training Accuracy:',logreg.score(X_train,y_train))
print('Testing Accuracy:',logreg.score(X_test,y_test))
```

Training Accuracy: 0.8046144816961338

Testing Accuracy: 0.8050326461741618

1. Model accuracy is 80%. Which can be further improved.
2. From coefficients, sub_grade and interest rate are the features which impart most effects in predicting the loan status.
3. Managing unbalancing in the data may help to improve the model.

2.Display model coefficients with column names

```
[ ]: print(list(zip(X_train.columns,np.round(logreg.coef_,4)[0])))

[('loan_amnt', 0.0195), ('term', -0.1907), ('int_rate', 0.6189), ('installment',
-0.0828), ('grade', 0.0439), ('sub_grade', -1.1288), ('emp_length', 0.03),
('home_ownership', -0.1257), ('annual_inc', 0.1548), ('verification_status',
-0.026), ('dti', -0.1992), ('open_acc', -0.1123), ('revol_bal', 0.0693),
('revol_util', -0.0926), ('total_acc', 0.1145), ('initial_list_status', 0.0334),
('application_type', 0.0152), ('pub_rec_flag', -0.0221), ('mort_acc_flag',
0.0579), ('pub_rec_bankruptcies_flag', 0.0035)]
```


###4.Results Evaluation

1.ROC AUC Curve & comments

```
[ ]: #importing roc
from sklearn.metrics import roc_curve, roc_auc_score
```

```
[ ]: prob = logreg.predict_proba(X_test)
```

```
[ ]: #roc curve
probabilites = prob[:,1]
fpr, tpr, thr = roc_curve(y_test,probabilites)
plt.plot(fpr,tpr)
plt.title('ROC curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.show()
```



```
[ ]: #auc of roc curve
roc_auc_score(y_test,probabilites)
```

```
[ ]: 0.7067709802132277
```

Comments

1. ROC curve shows a moderately good model performance.
2. Area under ROC curve is 0.7.
3. Since the data is imbalanced, ROC curve might be dominated by the majority class.


2.Precision Recall Curve & comments

```
[ ]: #importing pr curve
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import auc

precision, recall, thr = precision_recall_curve(y_test, probabilites)
plt.plot(recall, precision)

plt.xlabel('Recall')
```

```
plt.ylabel('Precision')
plt.title('PR curve')
plt.show()
```



```
[ ]: #auc of pr curve
auc(recall, precision)
```

```
[ ]: 0.8985007948437684
```

Comments

1. Precision-Recall curve shows good model performance.
2. The area under PR curve is 0.898.
3. Classification Report (Confusion Matrix etc)

```
[ ]: from sklearn.metrics import f1_score
train_y_pred = logreg.predict(X_train)
test_y_pred = logreg.predict(X_test)

train_score = f1_score(y_train, train_y_pred)
```

```
test_score = f1_score(y_test, test_y_pred)

print('Training F1 Score:', train_score)
print('Testing F1 Score:', test_score)
```

Training F1 Score: 0.8893629218524118

Testing F1 Score: 0.8897005113221329


Confusion Matrix

```
[ ]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

conf_matrix = confusion_matrix(y_test, test_y_pred)

ConfusionMatrixDisplay(conf_matrix).plot()
```

```
[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7de79ead5390>
```



###Tradeoff Questions:

How can we make sure that our model can detect real defaulters and there are less false positives?
This is important as we can lose out on an opportunity to finance more individuals and earn interest

on it.

The goal here is to reduce real defaulters i.e, reduce false positives. A **high precision** model can improve the performance in this aspect. 1. In hyper parameter tuning, there are methods to choose hyper-parameter which will ensure higher precision like precision scorer. 2. Class weights can be used to penalize errors in one class more than errors in the other.

Since NPA (non-performing asset) is a real problem in this industry, it's important we play safe and shouldn't disburse loans to anyone.

Non-performing assets are loans of which repayment is not happening for a long period of time. Avoiding NPAs are important in maintaining healthy functioning of the firm. Here also avoiding false positives is the goal.

Actionable Insights & Recommendations

- Model accuracy is 80%. Which can be further improved.
- From coefficients, sub_grade and interest rate are the features which impart most effects in predicting the loan status.
- Managing unbalancing in the data may help to improve the model.
- ROC curve shows a moderately good model performance. Area under ROC curve is 0.7. Since the data is imbalanced, ROC curve might be dominated by the majority class.
- Precision-Recall curve shows good model performance. The area under PR curve is 0.898.
- Non-performing assets are loans of which repayment is not happening for a long period of time. Avoiding NPAs are important in maintaining healthy functioning of the firm.
- The goal here is to reduce real defaulters i.e, reduce false positives. A high precision model can improve the performance in this aspect.

In hyper parameter tuning, there are methods to choose hyper-parameter which will ensure higher precision like precision scorer. Class weights can be used to penalize errors in one class more than errors in the other.
